A Quantitative Synthesis of Early Language Acquisition Using Meta-Analysis

Molly Lewis[1], Mika Braginsky[1], Sho Tsuji[2], Christina Bergmann[2], Page Piccinini[2],

Alejandrina Cristia[2], Michael C. Frank[1]

[1] Department Psychology, Stanford University

[2] Laboratoire de Sciences Cognitives et Psycholinguistique, ENS

Author note

Correspondence concerning this article should be addressed to Molly Lewis, Psychology

Department, Stanford University. 450 Serra Mall, Stanford, CA 94305. E-mail:

mll@stanford.edu.

Abstract

replicability, etc.

Keywords: replicability, reproducibility, meta-analysis, developmental psychology, language acquisition

Word count: XXXX

A Quantitative Synthesis of Early Language Acquisition Using Meta-Analysis

Introduction

To learn to speak a language, a child must acquire a wide range knowledge and skills: the sounds of the language, the word forms, and how the word forms map onto meanings, to name only a few. How does this process unfold? Our goal as psychologists is to build a theory that can explain and predict these behaviors in a way that is both precise but also highly generalizable. The challenge we face, however is that we must build this theory on the basis of very limited data. And, as a consequence, we must rely on inductive reasoning strategies to go from the data we observe to the theories we posit. In this paper, we demonstrate how meta-analytic methods can support the inductive process of theory building in the case of language acquistion.

Meta-analysis is a quantiatiave method for aggregating across experimental findings. The fundamental unit of meta-analysis is the *effect size*: a scale-free, quantitative measure of "success" in a phenomenon. Importantly, an effect size provides an estimate of the *size* of an effect, as well as a measure of uncertainty around this point estimate. With such a quantitative measure of success, we can apply the same reasoning we use to aggregate noisy measurements over participants in a single study: By assuming each *study*, rather than participant, is sampled from a population, we can appeal to the classical statistical framework to combine estimates of the effect size for a given phenomenon.

Meta-analytic methods support theory building in several ways. First, they provide a way of evaluating which effects in a literature are "real," and thus should constrain the theory. The field of psychology has grown concerned recently in the face of high-profile evidence that an effect observed in one study may not be the same in another ("replicability crisis"; Ioannidis, 2005; Nosek, 2012, 2015). Importantly, however, some variability across studies is to be expected due to sampling error alone. Effect sizes and meta-analytic methods provide a principled approach for distinguising variability due to sampling error from variability in the underyling effect. With a more veridical description of the empirical

landscape, we are in a better position to build a theory.

Second, meta-analysis support theory building by providing higher fidelity descriptions of phenomenona. Rather than simply concluding that an effect exists, effect sizes allow us to describe as, for exapmple, How does the effect change over development? How does the effect compare in magnitude to other effects? How much does a moderator affect the effect? This type of continuous analysis supports building quantitative models, and specifying theories that are more precise and constraining.

(SYNTHESIS POINT) Furthermore, to build parsimonious theories of human behavior, we should seek to explain not just individual phenemenon, but entire literatures of research. What is needed, then, is a tool for aggregating noisy data across studies within a phenomenon, as well as a common language for comparing effects across phenomenona.

Finally, in addition to these theoretical motivations, there are practical reasons for conducting a quantitative synthesis. When planning an experiment, an estimate of the size of an effect on the basis of prior literature can inform the sample size needed to achieve a desired level of power. Meta-analytic estimates of effect sizes can also aid in design choices: If a certain paradigm tends to have overall larger effect sizes than another, the strategic researcher might select this paradigm in order to maximize the power of a study.

While meta-analytic methods are likely helpful for many psychological literatures, we beleive language acquisition is a particularly informative application for this tool. One reason is that language acquistion may be uniquely vulnerable to false findings because running children is expensive, and thus sample sizes are small and studies are underpowered (Ioanndis, 2005). In addition, the difficulty in running participants means that replications are relatively rare in the field. Finally, there has been recent attention to developmental research practices more broadly, and corn (Peterson, 2016).

Towards this end, we have developed a dataset of effect sizes in the language acquistion literature, across 12 phenomenona, and XXX studies. Metalab. We have two goals:

- Describe the state of the field in terms of its participation in practices that are

prerequisites to cumulative science, and ultimately, a theoretical synthesis

- Provide a preliminary theoretical synthesis of the field

Veridical

Fidelity * across development * across different theoretical moderators (does X matter?) * across phenomenoa

METALAB

This problem of induction is particularly challenging in the case of behaviors that rely on many different skills, such as langauge acquisition. To acquire a language

In achieving this goal, however, we are fundamentally challeged e challenge we face, however, is that we must build these theories on the basis of limited data, typically collected from only 20-30 participants. This means that we must infer com

The developmental trajectory of language aquistion is

Psychologists hope to build generalizable theories about human behavior—theories that hold true beyond particulars of an individual study.

This quantitative approach provides a rich tool kit for synthesizing across literatures. By describing different phenomena using the same unit of measurement, we are able to compare effects in different domains.

In practice, however, the feasability of this meta-analytic approach relies on the field's commitment to practices that facilitate cumulative science. These practices apply to all stages of the research process. At the stage of experimental planning, researchers must pre-specify analytical descision to limit "researcher" degrees of freedom (Simmons, 2011; Simonsohn, 2014a, 2014b, 2014c). At the stage of completion, researchers should share a result regardless of its significance (Rosenthal, 1979; Fanelli 2012). And, at the stage of sharing, researchers must provide enough information about the method for another lab to conduct a close replication. Critically,r eports must also contain complete descriptions of both data and analytical decisions so that effect sizes can be calculcated for the purposes of meta-analysis,

Method

We analyzed 12 different phenomenena in language acquisition. We selected these phenomena in order to describe development at many different levels of the language hierarchy, from the acquistion of prosody and phonemic contrasts, to gaze following in linguistic interaction. This wide range of phenomena allowed us to compare the course of development across different domains, as well as explore questions about the interactive nature of language acquisition (Table 1).

To obtain estimates of effect size, we coded papers reporting experimental data. Within each paper, we calculated a separate effect size estimate for each experiment and age group ("condition"). In total, our sample includes estimates from 269 papers, 981 different conditions and 11,493 participants. The process for selecting papers from the literature differed by domain, with some individual meta-analyses using more systematic approaches than others. [Simulations here?]

Replicability of the field

A literature is more likely to describe a real effect if studies are randomly sampled from the population of all possible studies that researchers could in principle conduct. This assumption does not mean, however, that there should be *no* variability in effect size across studies: We should expect random variation around the true mean effect size, with smaller studies showing more variability around this mean.

Variability in effect sizes will be biased when this assumption of random study sampling does not hold. Bias may be introduced by the experimenter in a number of ways, including failure to publish null findings (Fanelli, 2010; Rosenthal, 1979; "publication bias", Rothstein, Sutton, & Borenstein, 2006), analytical flexibility (e.g., "p-hacking," Simmons, Nelson, & Simonsohn, 2011; Simonsohn, Nelson, & Simmons, 2014), reporting errors, or even fraud. These biases are problematic for theoretical development because they lead to large but often unknown errors in estimates of the effect size. If bias is present in the literature,

Level	Phenomenon	Description	N papers (conditions)
Prosody	IDS preference (Dunst, Gorman, & Hamby, 2012)	Looking times as a function of whether infant-directed vs. adult-directed speech is presented as stimulation.	16 (50)
Sounds	Phonotactic learning (Cristia, in prep.)	Infants' ability to learn phonotactic generalizations from a short exposure.	15 (47)
	Vowel discrimination (native) (Tsuji & Cristia, 2014)	Discrimination of native-language vowels, including results from a variety of methods.	40 (167)
	Vowel discrimination (non-native) (Tsuji & Cristia, 2014)	Discrimination of non-native vowels, including results from a variety of methods.	21 (72)
	Statistical sound learning (Cristia, in prep.)	Infants' ability to learn sound categories from their acoustic distribution.	11 (40)
	Word segmentation (Bergmann & Cristia, 2015)	Recognition of familiarized words from running, natural speech using behavioral methods.	68 (296)
Words	Mutual exclusivity (Lewis & Frank, in prep.)	Mapping of novel words reflecting children's inference that novel words tend to refer to novel objects.	20 (60)
	Sound Symbolism (Lammertink et al., in prep.)	Relationships between form and meaning (the bouba-kiki effect).	10 (42)
	Concept-label advantage (Lewis & Long, unpublished)	Infants' categorization judgments in the presence and absence of labels.	16 (100)
	Online word recognition (Frank, Lewis, & MacDonald, 2016)	Online word recognition of familiar words using two-alternative forced choice preferential looking.	12 (32)
Communication	Gaze following (Frank, Lewis, & MacDonald, 2016)	Gaze following using standard multi-alternative forced-choice paradigms.	15 (45)
	Pointing and vocabulary (Colonnesi et al., 2010)	Longitudinal correlations between declarative pointing and later vocabulary.	25 (30)

Table 1

Overview of meta-analyses in dataset.

estimates of effect size may be poor estimates of the true underlying effect size and thus be of limited evidential value. To make theoretical progress, we must therefore distinguish variability in effect sizes due to sample size from varability due to bias.

To assess the replicability of language acquisition phenomena, we conducted several

diagnostic analyses: Meta-analytic estimates of effect size, fail-safe-N (Orwin, 1983), funnel plots, and p-curve (U. Simonsohn, Nelson, & Simmons, 2014; Simonsohn et al., 2014; Simonsohn, Simmons, & Nelson, 2015). These analytical approaches each have limitations, but taken together, they provide converging evidence about the replicability of a literature. Overall, we find little evidence of bias in our meta-analyses, suggesting that the language acqusition literature likely describes real psychological phenomenona and should therefore provide the basis for theoretical development.

Meta-Analytic Effect Size

Meta-analysis provides a quantitative method for aggregating across studies. To estimate the overall effect size of a literature, effect sizes are pooled across papers to obtain a single meta-analytic estimate. Importantly, meta-analysis allows us to model variability in effect sizes due to differences in sample sizes by weighting studies with more participants more heavily in the overall estimate. This meta-analytic effect-size can be thought of as the "best estimate" of the effect size for a phenomenon given all the available data in the literature.

Table 2, column 4 presents meta-analytic estimates for each of our phenomenona. We find evidence for a non-zero effect size in 11 out of 12 of our phenomena, suggesting these literature provide evidential value. In the case of phonotatic learning, however, we find that the meta-analytic effect size estimate does not differ from zero, suggest that this literature does not describe a real effect. [Remove it from analyses below?].

While the meaure of effect size is itself quantitative, meta-analytic estimates of effect size provide only categorical information about the evidential value of a literature: the effect is real, or not. But, a more powerful method of assessing evidential value would tell us the *degree* to which a literature has evidential value, and thus the degree to which it shuold constrain our theory building. In the following three analyses—fail-safe-N, funnel plots, and p-curves—we describe through analyses that quantify the evidential value of these literatures.

Fail-safe-N

One approach for quantifying the reliability of a literature is to ask, How many missing studies with null effects would have to exist in the "file drawer" in order for the overall effect size to be zero? This is called the "fail-safe" number of studies (Orwin, 1983). To answer this question, we estimated the overall effect size for each phenomenenon (Table 2, column 2), and then used this to estimate the fail-safe-N (Table 2, column 3).

This analysis suggests a very large number of studies would have to be "missing" in each literature ($M = 3634$) in order for the overall effect sizes to be 0. Thus, while it is possible that some reporting bias is present in the literature, the large fail-safe-N suggests that the literature nonetheless likely describes a real effect.

One limitation of this analysis, however, is that it assumes that all reported effect sizes are obtained in the absence of analytical flexibility: If experimenters are exercising analytical flexibility through practices like p-hacking, then then number and magnitude of observed true effects in the literature may be inflated. In the next analysis, we examine this possibility through funnel plots.

Funnel Plots

Funnel plots provide a visual method for evaluating whether variability in effect sizes is due only to differences in sample size. A funnel plot shows effect sizes versus a metric of sample size, standard error. If there is no bias in a literature, we should expect studies to be randomly sampled around the mean, with more variability for less precise studies.

Figure 1 presents funnel plots for each of our 12 meta-analyses. These plots show evidence of asymmetry (bias) for several of our phenomenon (Table 2, column 4). However, an important limitation of this method is that it is difficult to determine the source of this bias. One possibility is that this bias is due not to experimenter, but to true heterogenity in phenomena (e.g. different ages). P-curve analyses provide one method for addressing this issue, which we turn to next.
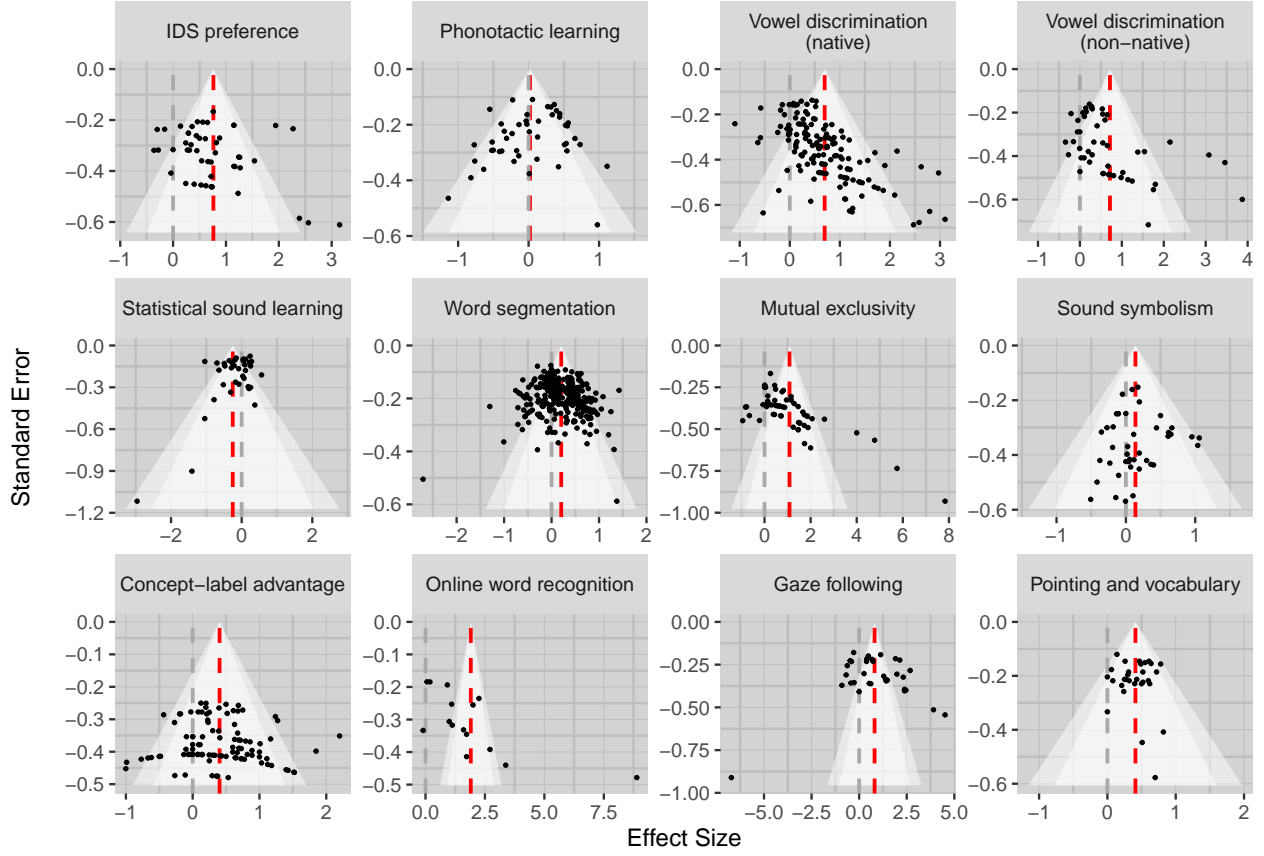
Figure 1. Funnel plots for each meta-analysis. Each effect size estimate is represented by a point, and the mean effect size is shown as a red dashed line. The funnel corresponds to a 95% (narrow) and 99% (wide) CI around this mean. In the absense of true heterogenity in effect sizes (no moderators) and bias, we should expect all points to fall inside the funnel.

P-curves

P-curves provide a more robust way to identify bias in a literature (U. Simonsohn et al., 2014; Simonsohn et al., 2014, 2015). A p-curve is the distribution of p-values for the statistical test of the main hypothesis across a literature. Critically, if there is a real effect in the literature, the shape of the p-curve should reflect this. In particular, we should expect the p-curve to be right skewed with more small values (e.g., .01) than large values (e.g., .04). An important features of this method is that we should expect this skew independent of any true heterogenity in the data, such as age. Evidence that the curve is in fact right-skewed would suggest that the literature is not biased, and that it provides evidential value for theory building.
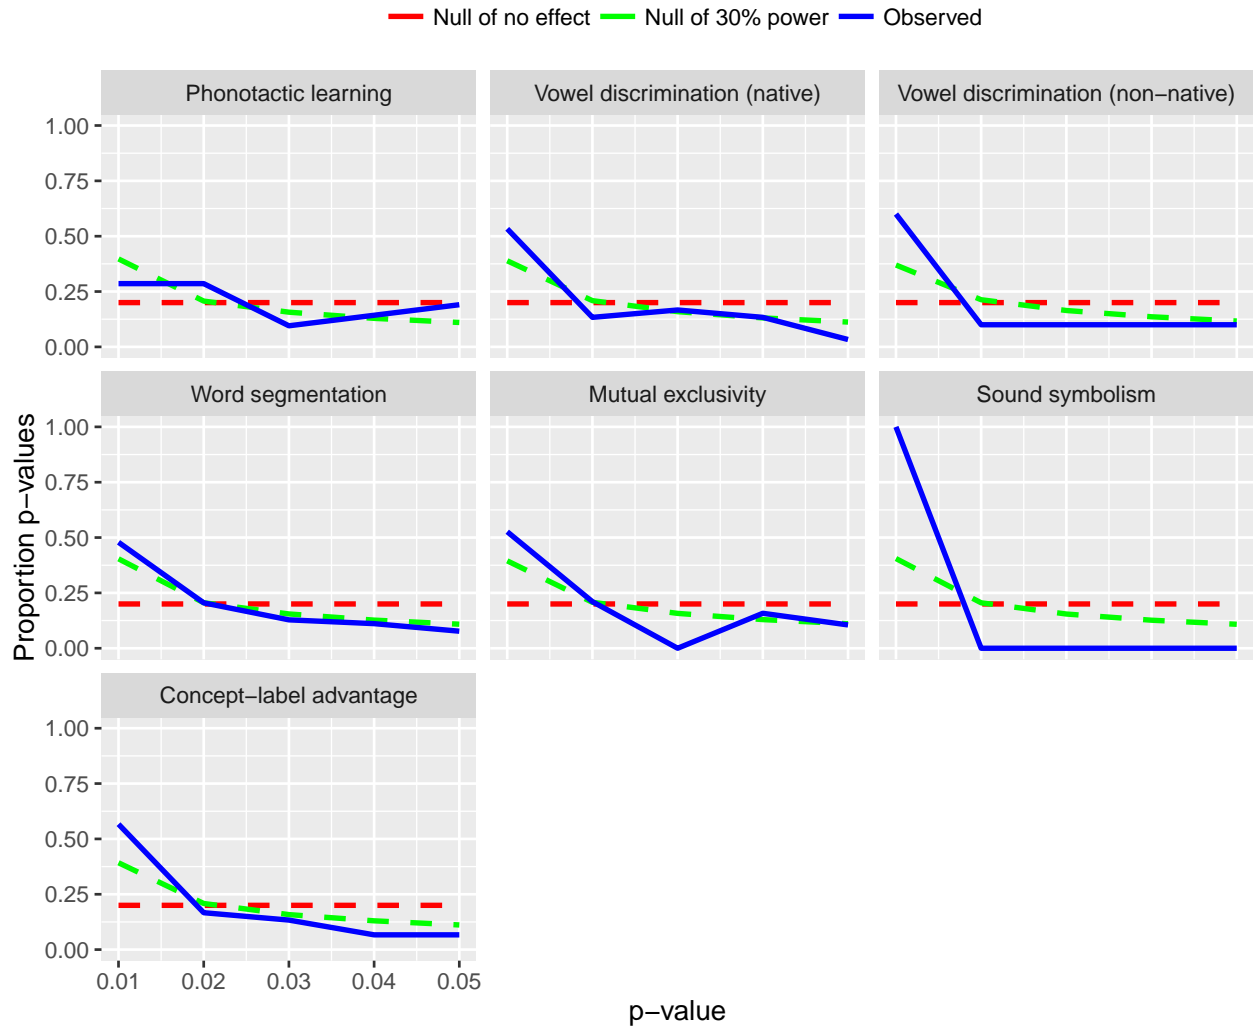
Figure 2. P-curve for each meta-analysis (Simonsohn, Nelson, & Simmons, 2014). In the absense of p-hacking, we should expect the observed p-curve (blue) to be right-skewed (more small values). The red dashed line shows the expected distribution of p-values when the effect is non-existent (the null is true). The green dashed line shows the expected distribution if the effect is real, but studies only have 33% power.

Phenomenon	*d*	fail-safe-N	funnel skew	p-curve skew	power
IDS preference	0.71 [0.53, 0.89]	3762	1.88 (0.06)		
Phonotactic learning	0.04 [-0.09, 0.16]	45	-1.08 (0.28)	-1.52 (0.06)	0.14
Vowel discrimination (native)	0.6 [0.5, 0.71]	9536	8.98 (0)	-5.42 (0)	0.67
Vowel discrimination (non-native)	0.66 [0.42, 0.9]	3391	4.13 (0)	-3.24 (0)	0.78
Statistical sound learning	-0.14 [-0.27, -0.02]	Inf	-1.87 (0.06)		
Word segmentation	0.2 [0.15, 0.25]	5645	1.54 (0.12)	-9.67 (0)	0.56
Mutual exclusivity	1.01 [0.68, 1.33]	6443	6.25 (0)		
Sound symbolism	0.15 [0.04, 0.26]	538	-1.32 (0.19)	-2.16 (0.02)	0.96
Concept-label advantage	0.4 [0.29, 0.51]	3928	0.31 (0.76)	-6.15 (0)	0.69
Online word recognition	1.89 [0.81, 2.96]	2843	2.92 (0)		
Gaze following	0.84 [0.26, 1.42]	2641	-1.69 (0.09)		
Pointing and vocabulary	0.41 [0.32, 0.49]	1202	0.59 (0.55)		

Table 2

Summary of replicability analyses. d = Effect size (Cohen's d) estimated from a random-effect model; fail-safe-N = number of missing studies that would have to exist in order for the overall effect size to be d = 0; funnel skew = test of asymmetry in funnel plot using the random-effect Egger's test (Stern & Eggers, 2005); p-curve skew = test of the right skew of the p-curve using the Stouffer method (Simonsohn, Simmons, & Nelson, 2015); power = power to reject the null hypothesis at the 5% significance level based on the p-curve (Simonsohn, Nelson, & Simmons, 2014); Brackets give 95% confidence intervals, and parentheses show p-values.

Figure 2 shows p-curves for 7 of our 12 meta-analyses[1]. Across all p-curves, the curves show evidence of right skew. This is confirmed by formal analyses (Table 2, column 5). P-curves also provide a method for calculcating the overall power of a literature, based on the shape of the p-curve (U. Simonsohn et al., 2014). Intuitively, when power is high and effect is real, we should be more likely to observe an effect size "further" from the null. This means that we will observe more small effect sizes. Thus, the higher the power, the more right skewed the p-curve will be. Table 2 (column 6) presents estimates of power for each meta-analysis based on p-curve. With the exception of phonotactic learning (*power* = .14),

---

[1]We were unable to do this analysis on all meta-analyses because some were missing key statistical tests (e.g. gaze following) or the test statistic was not available (e.g. pointing and vocabulary).

literatures appear to have acceptable power.

Theoretical Synthesis

Given this literature has evidential value, we next turn to drawing inferences on the basis of the literature.

Statistical Approach

METAMETAPLOT

Discussion

Limitations

Author Contributions.

Acknowledgments.

References. Bergmann, C., & Cristia, A. (2015). Development of infants' segmentation of words from native speech: A meta-analytic approach. *Developmental Science.*

Dunst, C., Gorman, E., & Hamby, D. (2012). Preference for infant-directed speech in preverbal young children. *Center for Early Literacy Learning, 5*(1).

Fanelli, D. (2010). Positive Results Increase Down the Hierarchy of the Sciences. *PLoS ONE, 5*(4), e10068–10.

Frank, M. C., Lewis, M. L., & MacDonald, K. (in press). A performance model for early word learning. In *Proceedings of the 38th Annual Conference of the Cognitive Science Society.* Retrieved from

http://langcog.stanford.edu/papers_new/frank-2016-underrev.pdf

Lewis, M., & Frank, M. C. (in prep). Multiple routes to disambiguation.

Orwin, R. G. (1983). A fail-safe n for effect size in meta-analysis. *Journal of Educational Statistics*, 157–159.

Peterson, D. (2016). The Baby Factory: Difficult Research Objects, Disciplinary Standards, and the Production of Statistical Significance. *Socius: Sociological Research for a Dynamic World, 2*(0), 1–10.

Rosenthal, R. (1979). The file drawer problem and tolerance for null results. *Psychological Bulletin, 86*(3), 638.

Rothstein, H. R., Sutton, A. J., & Borenstein, M. (2006). *Publication bias in meta-analysis: Prevention, assessment and adjustments.* John Wiley & Sons.

Simmons, Nelson, L. D., & Simonsohn, U. (2011). False-Positive Psychology: Undisclosed Flexibility in Data Collection and Analysis Allows Presenting Anything as Significant. *Psychological Science, 22*(11), 1359–1366.

Simonsohn, U., Nelson, L. D., & Simmons, J. P. (2014). P-curve and effect size correcting for

publication bias using only significant results. *Perspectives on Psychological Science,*
9(6), 666–681.

Simonsohn, Nelson, L. D., & Simmons, J. P. (2014). P-curve: A key to the file-drawer.
Journal of Experimental Psychology: General, 143(2), 534.

Simonsohn, Simmons, J. P., & Nelson, L. D. (2015). Better p-curves. *Simonsohn, Uri,*
Joseph P. Simmons, and Leif D. Nelson (Forthcoming),"Better P-Curves," Journal of
Experimental Psychology: General.

Sterne, J. A., & Egger, M. (2005). Regression methods to detect publication and other bias
in meta-analysis. *Publication Bias in Meta-Analysis: Prevention, Assessment, and*
Adjustments, 99–110.

Tsuji, S., & Cristia, A. (2014). Perceptual attunement in vowels: A meta-analysis.
Developmental Psychobiology, 56(2), 179–191.